

Why MongoDB Is Awesome

DevNation Chicago, IL
May 15, 2010

John Nunemaker
Ordered List



I am user #4,243 on Twitter

@jnunemaker



**...the best features of key/
values stores, document
databases and relational
databases in one.**

John Nunemaker

RailsTips.org June '09

Created by    

10gen

I



mongoDB

I



A LOT



mongoDB

Which has led some people to believe

that I am on the payroll.



...but I am not.

I am merely a

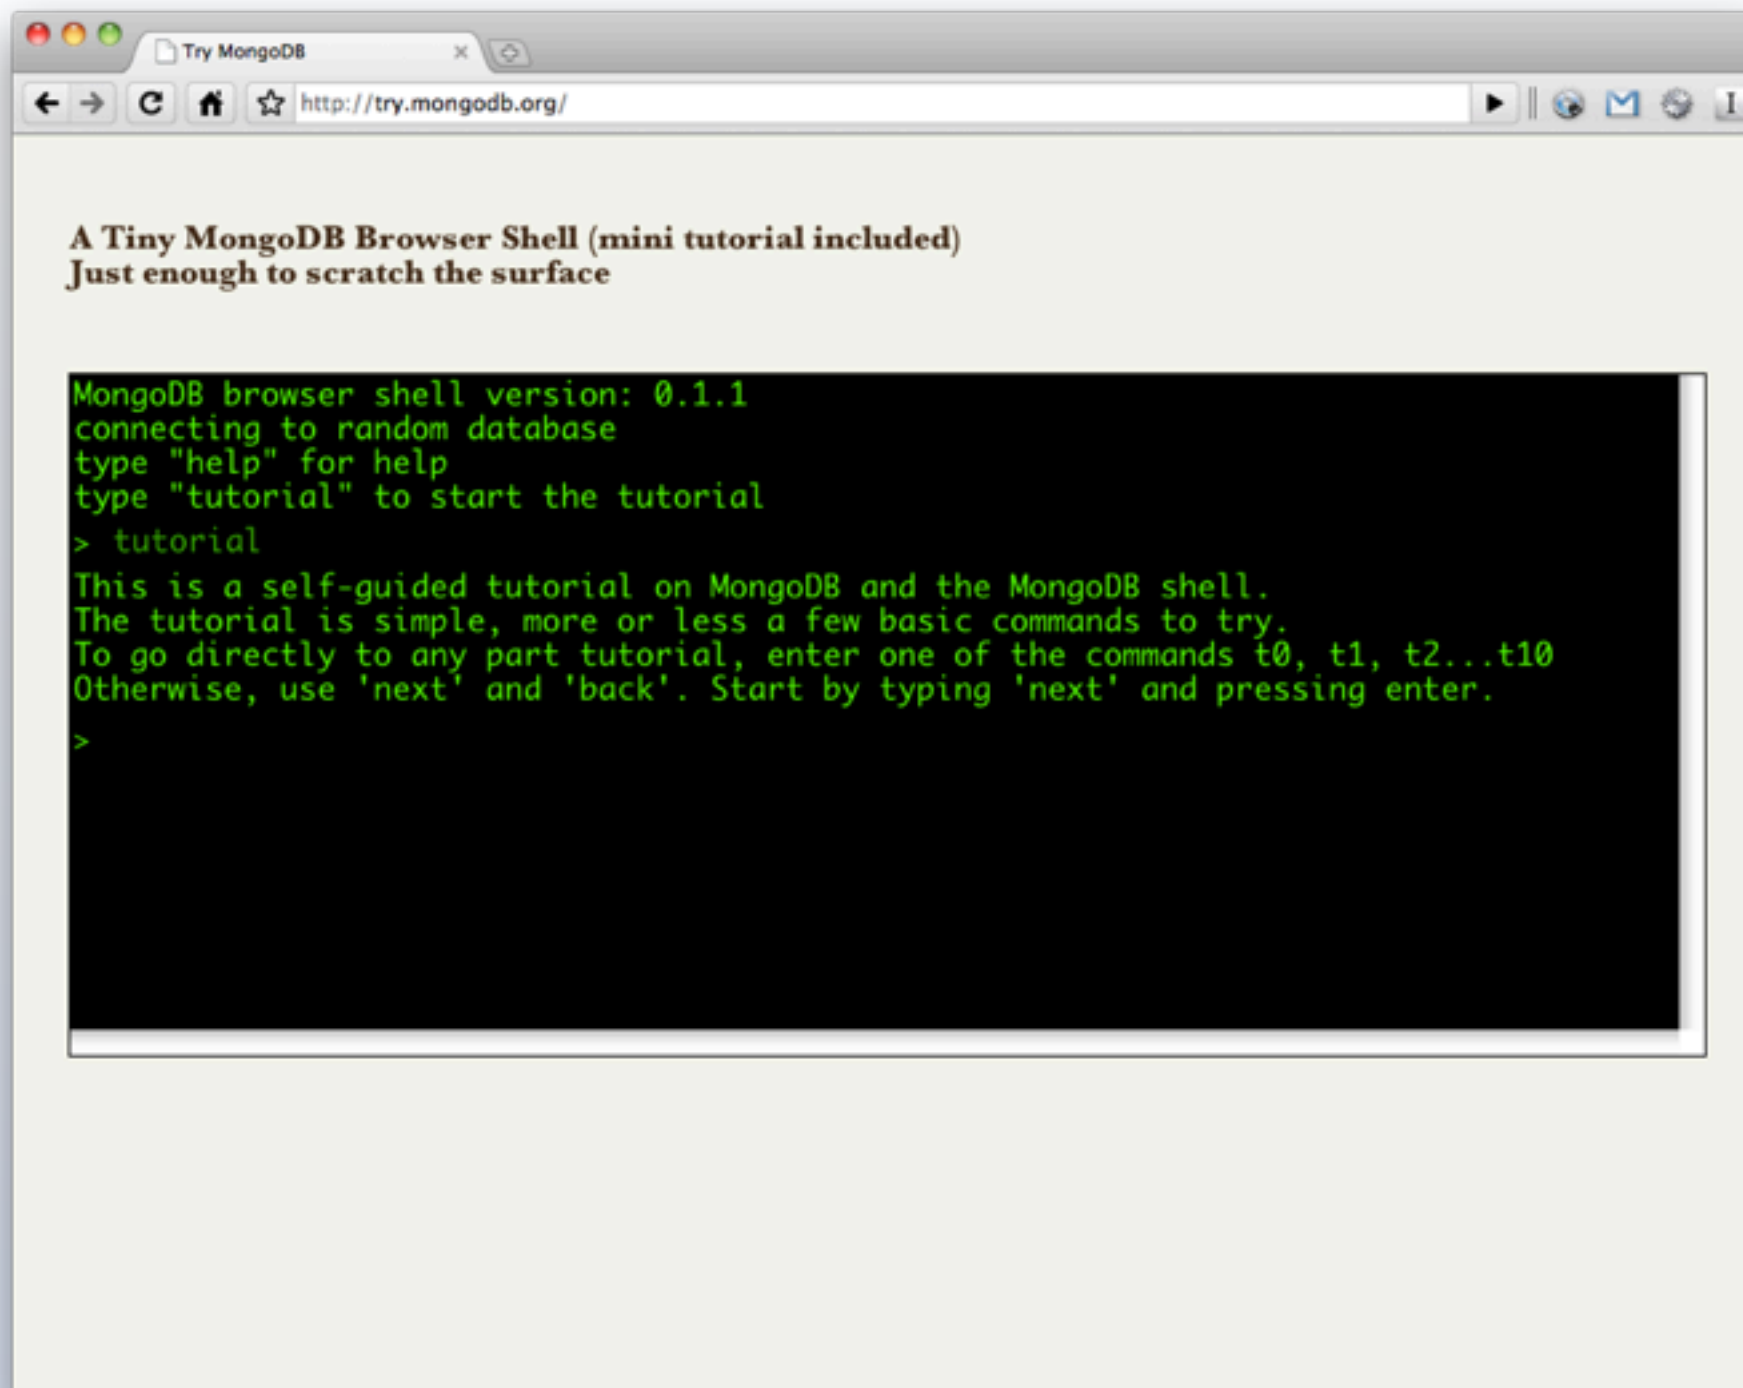
Satisfied User

Easy To

Try

Easy to Try

In Your Browser



<http://try.mongodb.org/>

Easy to Try

On Your Computer

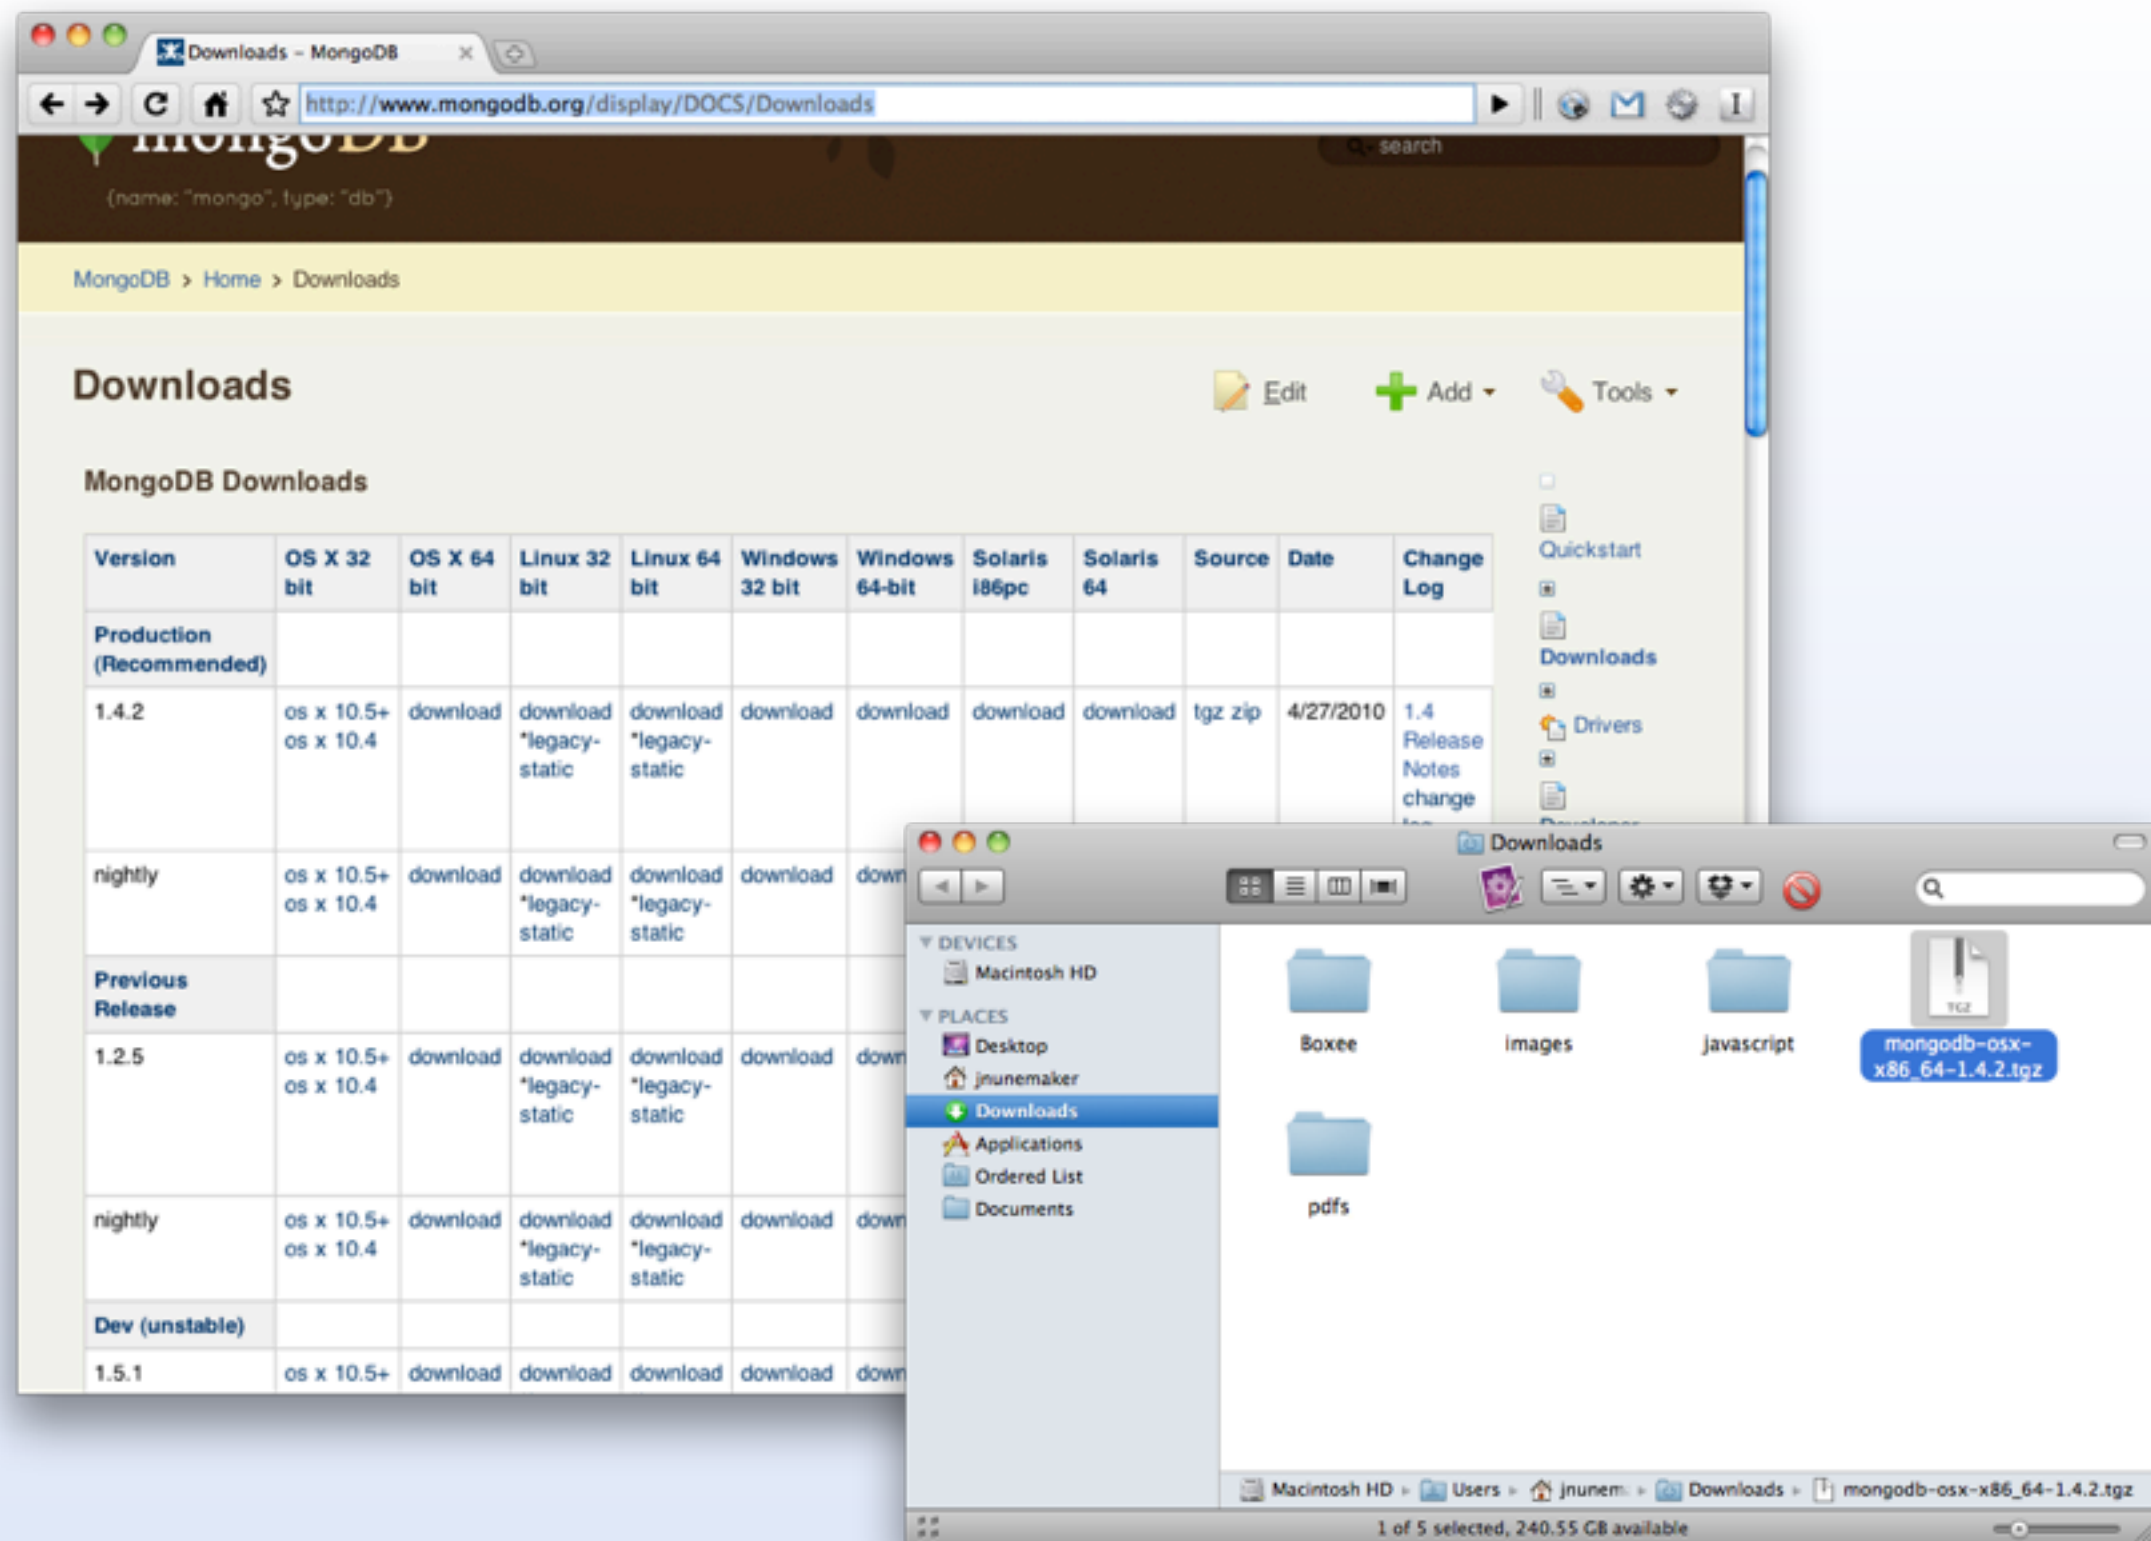
```
$ wget http://downloads.mongodb.org/osx/mongodb-osx-x86_64-1.4.2.tgz
```

```
$ wget http://downloads.mongodb.org/osx/mongodb-osx-x86_64-1.4.2.tgz  
$ tar -xf mongodb-osx-x86_64-1.4.2.tgz
```

```
$ wget http://downloads.mongodb.org/osx/mongodb-osx-x86_64-1.4.2.tgz
$ tar -xf mongodb-osx-x86_64-1.4.2.tgz
$ mkdir -p /data/db
```




```
$ wget http://downloads.mongodb.org/osx/mongodb-osx-x86_64-1.4.2.tgz
$ tar -xf mongodb-osx-x86_64-1.4.2.tgz
$ mkdir -p /data/db
$ mongodb-osx-x86_64-1.4.2/bin/mongod
```



<http://www.mongodb.org/display/DOCS/Downloads>

Easy to Try

From Your Language



<http://www.mongodb.org/display/DOCS/Drivers>

Easy To

Understand

Easy to Understand

Similar Terms

Database == Database

- > show dbs
admin
harmony-development
harmony-test
local
...
- > use harmony-development
switched to db harmony-development
- > show collections
accounts
activities
assets
items
...

Collection == Table

```
> db.accounts  
harmony-development.accounts
```

```
> db.accounts.count()  
1
```

```
> db.accounts.find().forEach(function(doc) {  
    print(tojson(doc));  
});
```

Document == Row

```
{
  "_id"      : ObjectId("4be97eaebcd1b30e86000003"),
  "title"    : "Ordered List",
  "creator_id" : ObjectId("4be97eadbcd1b30e86000001"),
  "memberships" : [
    ObjectId("4be97eadbcd1b30e86000001"),
    ObjectId("4be97eaebcd1b30e86000002")
  ]
}
```

Easy to Understand

Similar Functionality

Dynamic Queries

<http://www.mongodb.org/display/DOCS/Querying>

<http://www.mongodb.org/display/DOCS/Advanced+Queries>

> use testing
switched to db testing

```
> db.colors.insert({name: 'red',      primary: true})
> db.colors.insert({name: 'green',    primary: true})
> db.colors.insert({name: 'blue',     primary: true})
> db.colors.insert({name: 'purple',   primary: false})
> db.colors.insert({name: 'orange',   primary: false})
> db.colors.insert({name: 'yellow',   primary: false})
```

```
> var cursor = db.colors.find()
> cursor.next()
{
  "_id" : ObjectId("4bed7aeb0b4acd070c593ba6"),
  "name" : "red",
  "primary" : true
}
```


> cursor

```
{ "_id" : ObjectId("4bed7af40b4acd070c593ba7"), "name" : "green", "primary" : true }
{ "_id" : ObjectId("4bed7af80b4acd070c593ba8"), "name" : "blue", "primary" : true }
{ "_id" : ObjectId("4bed7b570b4acd070c593ba9"), "name" : "purple", "primary" : false }
{ "_id" : ObjectId("4bed7b6a0b4acd070c593baa"), "name" : "orange", "primary" : false }
{ "_id" : ObjectId("4bed7b7d0b4acd070c593bab"), "name" : "yellow", "primary" : false }
```

```
SELECT * from colors WHERE name = 'green'
```

```
SELECT * from colors WHERE name = 'green'
```

```
> db.colors.find({name: 'green'})
```

```
{ "_id" : ObjectId("4bed7af40b4acd070c593ba7"), "name" : "green", "primary" : true }
```

```
SELECT name from colors WHERE primary = 1
```

```
SELECT name from colors WHERE primary = 1
```

```
> db.colors.find({primary:true}, {name:true})  
{ "_id" : ObjectId("4bed7aeb0b4acd070c593ba6"), "name" : "red" }  
{ "_id" : ObjectId("4bed7af40b4acd070c593ba7"), "name" : "green" }  
{ "_id" : ObjectId("4bed7af80b4acd070c593ba8"), "name" : "blue" }
```

```
> db.colors.find({name:/l/})
```

```
{ "_id" : ObjectId("4bed7af80b4acd070c593ba8"), "name" : "blue", "primary" : true }  
{ "_id" : ObjectId("4bed7b570b4acd070c593ba9"), "name" : "purple", "primary" : false }  
{ "_id" : ObjectId("4bed7b7d0b4acd070c593bab"), "name" : "yellow", "primary" : false }
```

```
> db.colors.find({primary:true}).sort({name:1}).limit(1)
{ "_id" : ObjectId("4bed7af80b4acd070c593ba8"), "name" : "blue", "primary" : true }
```

```
> db.colors.find({primary:true}).sort({name:1}).limit(1)
{ "_id" : ObjectId("4bed7af80b4acd070c593ba8"), "name" : "blue", "primary" : true }
```

```
> db.colors.find({primary:true}).sort({name:-1}).limit(1)
{ "_id" : ObjectId("4bed7aeb0b4acd070c593ba6"), "name" : "red", "primary" : true }
```



```
> db.colors.find({primary:true}).sort({name:1}).limit(1)
{ "_id" : ObjectId("4bed7af80b4acd070c593ba8"), "name" : "blue", "primary" : true }
```

```
> db.colors.find({primary:true}).sort({name:-1}).limit(1)
{ "_id" : ObjectId("4bed7aeb0b4acd070c593ba6"), "name" : "red", "primary" : true }
```

```
> db.colors.find({primary:true}).sort({name:1}).skip(1).limit(1)
{ "_id" : ObjectId("4bed7af40b4acd070c593ba7"), "name" : "green", "primary" : true }
```

```
> db.people.insert({name: 'John', age: 28})  
> db.people.insert({name: 'Steve', age: 29})  
> db.people.insert({name: 'Steph', age: 27})
```

```
SELECT * from people WHERE age > 27
```

```
SELECT * from people WHERE age > 27
```

```
> db.people.find({age: {$gt: 27}})
```

```
{ "_id" : ObjectId("4bed80b20b4acd070c593bac"), "name" : "John", "age" : 28 }
```

```
{ "_id" : ObjectId("4bed80bb0b4acd070c593bad"), "name" : "Steve", "age" : 29 }
```

```
SELECT * from people WHERE age <= 27
```

```
SELECT * from people WHERE age <= 27
```

```
> db.people.find({age: {$lte: 27}})
```

```
{ "_id" : ObjectId("4bed80c10b4acd070c593bae"), "name" : "Steph", "age" : 27 }
```

\$gt

\$gte

\$lt

\$lte

\$ne

\$in

\$nin

\$mod

\$all

\$size

\$exists

\$type

\$elemMatch

\$not

\$where

Indexes

<http://www.mongodb.org/display/DOCS/Indexes>


```
// single ascending
```

```
> db.colors.ensureIndex({name: 1})
```

```
// single ascending
```

```
> db.colors.ensureIndex({name: 1})
```

```
// single descending
```

```
> db.colors.ensureIndex({created_at: -1})
```

```
// single ascending
```

```
> db.colors.ensureIndex({name: 1})
```

```
// single descending
```

```
> db.colors.ensureIndex({created_at: -1})
```

```
// unique
```

```
> db.colors.ensureIndex({email: 1}, {unique: true})
```

```
// single ascending
```

```
> db.colors.ensureIndex({name: 1})
```

```
// single descending
```

```
> db.colors.ensureIndex({created_at: -1})
```

```
// unique
```

```
> db.colors.ensureIndex({email: 1}, {unique: true})
```

```
// non-blocking in background
```

```
> db.colors.ensureIndex({name: 1}, {background: true})
```

```
// single ascending
```

```
> db.colors.ensureIndex({name: 1})
```

```
// single descending
```

```
> db.colors.ensureIndex({created_at: -1})
```

```
// unique
```

```
> db.colors.ensureIndex({email: 1}, {unique: true})
```

```
// non-blocking in background
```

```
> db.colors.ensureIndex({name: 1}, {background: true})
```

```
// compound
```

```
> db.colors.ensureIndex({name: 1, created_at: -1})
```

Aggregation

<http://www.mongodb.org/display/DOCS/Aggregation>

```
> db.colors.count()
```

```
6
```

```
> db.colors.count  
({primary:true})
```

```
3
```

```
> db.colors.distinct('name')  
[ "blue", "green", "orange", "purple", "red", "yellow" ]
```

```
> db.people.distinct('name', {age:28})  
[ "John" ]
```



```
> db.items.insert({title:'Home',      template:'home'})
> db.items.insert({title:'What We Do', template:'page'})
> db.items.insert({title:'Our Writing', template:'page'})
> db.items.insert({title:'Who We Are',  template:'page'})
> db.items.insert({title:'Hire Us',     template:'page'})

> var key      = {template: true};
> var initial  = {count:0};
> var reduce   = function(obj, prev) { prev.count += 1; };

> db.items.group({key:key, initial:initial, reduce:reduce})
[
  {"template" : "home", "count" : 1},
  {"template" : "page", "count" : 4}
]
```

```
> db.items.insert({tags: ['dog', 'cat']})
> db.items.insert({tags: ['dog']})
> db.items.insert({tags: ['dog', 'mouse']})
> db.items.insert({tags: ['dog', 'mouse', 'hippo']})
> db.items.insert({tags: ['dog', 'mouse', 'hippo']})
> db.items.insert({tags: ['dog', 'hippo']})
```

```
> var map = function() {  
    this.tags.forEach(function(t) {  
        emit(t, {count: 1});  
    });  
}
```

```
> var reduce = function(key, values) {  
    var count = 0;  
    for(var i=0, len=values.length; i<len; i++) {  
        count += values[i].count;  
    }  
    return {count: count};  
}
```

```
> var result = db.items.mapReduce(map, reduce);
```

```
> var result = db.items.mapReduce(map, reduce);
> result
{
  "ok" : 1,
  "timeMillis" : 86,
  "result" : "tmp.mr.mapreduce_1273861517_683",
  "counts" : {
    "input" : 6,
    "emit" : 13,
    "output" : 4
  }
}
```

```
> db[result.result].find()
```

```
{ "_id" : "cat",    "value" : { "count" : 1 } }  
{ "_id" : "dog",    "value" : { "count" : 6 } }  
{ "_id" : "hippo",  "value" : { "count" : 3 } }  
{ "_id" : "mouse",  "value" : { "count" : 3 } }
```

Easy to Understand

Similar Data Types

**Array, Binary, Boolean, DateTime,
DB Reference, Embedded Object,
Integer, Null, ObjectId, RegExp,
String, Symbol, Timestamp**



BSON { 01010100 11101011 10101110 01010101 }

BSON [*bee · sahn*], short for Binary [JSON](#), is a binary-encoded serialization of JSON-like documents. Like JSON, BSON supports the embedding of documents and arrays within other documents and arrays. BSON also contains extensions that allow representation of data types that are not part of the JSON spec. For example, BSON has a Date type and a BinData type.

BSON can be compared to binary interchange formats, like [Protocol Buffers](#). BSON is more "schema-less" than Protocol Buffers, which can give it an advantage in flexibility but also a slight disadvantage in space efficiency (BSON has overhead for field names within the serialized data).

BSON was designed to have the following three characteristics:

1. Lightweight

Keeping spatial overhead to a minimum is important for any data representation format, especially when used over the network.

2. Traversable

BSON is designed to be traversed easily. This is a vital property in its role as the primary data representation for [MongoDB](#).

3. Efficient

Encoding data to BSON and decoding from BSON can be performed very quickly in most languages due to the use of C data types.

[specification](#)[implementation](#)[discussion](#)

```
> db.people.insert({
  name      : 'John',
  awesome   : true,
  shows     : ['Dexter', 'LOST', 'How I Met Your Mother'],
  info     : {
    age     : 28,
    home    : 'South Bend, IN',
    dob     : (new Date('November 25, 1981'))
  }
})
```

```
> var me = db.people.findOne({name: 'John'})
```

```
> me.name
```

```
John
```

```
> me.awesome
```

```
true
```

```
> me.shows[1]
```

```
LOST
```

```
> me.info.age
```

```
28
```

```
> me.info.dob.getFullYear()
```

```
1981
```

```
> db.people.find({'info.age': 28})  
{ "_id" : ObjectId("4bed9cba0b4acd070c593bc5"), "name" : "John" }
```

```
> db.people.find({'info.age': 28})  
{ "_id" : ObjectId("4bed9cba0b4acd070c593bc5"), "name" : "John" }
```

```
> db.people.find({shows:'Dexter'})  
{ "_id" : ObjectId("4bed9cba0b4acd070c593bc5"), "name" : "John" }
```

```
> db.people.find({'info.age': 28})
{ "_id" : ObjectId("4bed9cba0b4acd070c593bc5"), "name" : "John" }

> db.people.find({shows:'Dexter'})
{ "_id" : ObjectId("4bed9cba0b4acd070c593bc5"), "name" : "John" }

> db.people.find({shows:{$in:['Dexter', 'LOST']}})
{ "_id" : ObjectId("4bed9cba0b4acd070c593bc5"), "name" : "John" }
```

Easy to Understand

Similar Relationships

One to *Many*

1. Normalized

```
// insert post  
> db.posts.insert({title: 'Why Mongo Rocks'});  
> var post = db.posts.findOne({title: 'Why Mongo Rocks'});
```

```
// insert post
```

```
> db.posts.insert({title: 'Why Mongo Rocks'});
```

```
> var post = db.posts.findOne({title: 'Why Mongo Rocks'});
```

```
// insert comment
```

```
> db.comments.insert({  
  name      : 'John',  
  body      : 'Because...',  
  post_id   : post._id  
});
```

```
> var comment = db.comments.findOne({name: 'John'});
```

```
SELECT * FROM comments WHERE post_id = #{post.id}
```

```
> db.comments.find({post_id: post._id})  
{  
  "_id"      : ObjectId("4bee1cc79e89db4e12bf78de"),  
  "name"     : "John",  
  "body"     : "Because...",  
  "post_id"  : ObjectId("4bee1c519e89db4e12bf78dd")  
}
```

```
SELECT * FROM posts WHERE id = #{comment.id}
```

```
> db.posts.find({_id: comment.post_id})  
{  
  "_id"      : ObjectId("4bee1c519e89db4e12bf78dd"),  
  "title"    : "Why Mongo Rocks"  
}
```

2. Embedded

```
// insert post AND comments
> db.posts.insert({
  title: 'Why Mongo Rocks',
  comments: [
    {name: 'John', body: 'Because...'},
    {name: 'Steve', body: 'Uh huh!'}
  ]
})
```



```
> var post = db.posts.find({title: 'Why Mongo Rocks'});
```

```
> var post = db.posts.find({title: 'Why Mongo Rocks'});  
> post  
{  
  "_id"      : ObjectId("4bee21259e89db4e12bf78df"),  
  "title"    : "Why Mongo Rocks",  
  "comments" : [  
    {"name": "John", "body": "Because..."},  
    {"name": "Steve", "body": "Uh huh!"}  
  ]  
}
```

```
> db.posts.find({'comments.name': 'John'})
```

```
> db.posts.find({'comments.name': 'John'})
```

```
> db.posts.find({  
  comments: {  
    $elemMatch: {name: 'John'}  
  }  
})
```

```
// insert post AND comments AND threads!
> db.posts.insert({
  title: 'Why Mongo Rocks',
  comments: [
    {
      name: 'John',
      body: 'Because...',
      comments: [
        {name: 'Frank', body: 'You are crazy!'},
        {name: 'Billy', body: 'Frank Furter!'}
      ]
    }
  ]
})
```

```
> db.posts.insert({  
  title : 'Why Mongo Rocks',  
  tags  : ['mongodb', 'databases']  
})
```

```
> db.posts.insert({  
  title : 'Why Mongo Rocks',  
  tags  : ['mongodb', 'databases']  
})  
  
> db.posts.ensureIndex({tags:1})
```

Some Notes

Some Notes

- ❖ **Embedding is pre-joining**

Some Notes

- ❖ **Embedding is pre-joining**
- ❖ **Embed when document always appears with parent**

Some Notes

- ❖ **Embedding is pre-joining**
- ❖ **Embed when document always appears with parent**
- ❖ **4MB document size limit**

Many to Many

```
> db.sites.insert({domain: 'orderedlist.com'})
> db.sites.insert({domain: 'railstips.org'})
> db.sites.find()
{
  "_id"      : ObjectId("4bee280f9e89db4e12bf78e2"),
  "domain"   : "orderedlist.com"
}
{
  "_id"      : ObjectId("4bee283c9e89db4e12bf78e3"),
  "domain"   : "railstips.org"
}
```

```
> db.users.insert({  
  name: 'John',  
  authorizations: [  
    ObjectId('4bee280f9e89db4e12bf78e2'),  
    ObjectId('4bee283c9e89db4e12bf78e3')  
  ]  
})
```

```
> db.users.insert({  
  name: 'Steve',  
  authorizations: [  
    ObjectId('4bee280f9e89db4e12bf78e2')  
  ]  
})
```

```
> var orderedlist = db.sites.findOne({domain:'orderedlist.com'})  
> db.users.find({authorizations:orderedlist._id})  
// john and steve
```

```
> var railstips = db.sites.findOne({domain:'railstips.org'})  
> db.users.find({authorizations:railstips._id})  
// john
```

```
> var john = db.users.findOne({name: 'John'})  
> db.sites.find({_id: {$in: john.authorizations}})  
// orderedlist.com and railstips.org
```


Easy To

Learn

By Email

<http://groups.google.com/group/mongodb-user>

By IRC

`irc://irc.freenode.net/#mongodb`

By Web

<http://mongodb.org/>

<http://mongotips.com/>

By Book

<http://www.10gen.com/books>

<http://cookbook.mongodb.org/>

By Conference

<http://www.10gen.com/events>

<http://windycitydb.org/>

By Training

<http://ideafoundry.info/mongodb>

Thank you!

john@orderedlist.com

[@jnunemaker](#)

